

# Neural Networks: IV. Optimization (Part 3)

Jan Bauer

*jan.bauer@dhbw-mannheim.de*

16.05.19

- Find the  $W$  that minimizes the loss, i.e.  $W^* = \arg \min_W L$ 
  - Problem: Difficult. We only focus on the basics
- School: Gradient  $\stackrel{!}{=} 0$  & check Hessian  $\mathcal{H}$ 
  - **Example:**  $0 = \arg \min x^2$
  - Problem: Machine power intense

Find minima and maxima (recipe)

1  $\nabla_f(x) \stackrel{!}{=} 0$

2  $\mathcal{H}_f(x) \begin{cases} > 0 & \text{minimum} \\ < 0 & \text{maximum} \\ \text{else} & \text{saddlepoint} \end{cases}$

# Recap: Positive/negative definite

- A  $m \times m$  matrix  $M$  is positive definite if the leading principal minors are positive. We write  $M > 0$ .
- A  $m \times m$  matrix  $M$  is negative definite if the leading principal minors are negative. We write  $M < 0$ .
- Not relevant for this course, but for completion:
  - positive semi-definite if " $\geq 0$ ". We write  $M \geq 0$ .
  - negative semi-definite if " $\leq 0$ ". We write  $M \leq 0$ .

# Recap: Principle Minors (Example)

Let  $M = \begin{pmatrix} m_{1,1} & m_{1,2} & m_{1,3} \\ m_{2,1} & m_{2,2} & m_{2,3} \\ m_{3,1} & m_{3,2} & m_{3,3} \end{pmatrix}$ . The principal minors are then given by

- $M_{1,1} = \det(m_{1,1}) = m_{1,1}$
- $M_{2,2} = \det \begin{pmatrix} m_{1,1} & m_{1,2} \\ m_{2,1} & m_{2,2} \end{pmatrix} = m_{1,1} \cdot m_{2,2} - m_{1,2} \cdot m_{2,1}$
- $M_{3,3} = \det(M)$

# Recap: Gradient & Hessian

Let  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ ,  $x \mapsto f(x)$ . Then...

- $\nabla_f(x) \equiv \frac{\partial f}{\partial x} \equiv \sum_i \frac{\partial f}{\partial x_i} \mathbf{e}_i = \left( \frac{\partial f}{\partial x_1} \quad \cdots \quad \frac{\partial f}{\partial x_n} \right)'$
- $\mathcal{H}_f(x) \equiv \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{pmatrix}$

Find the extreme values of

- $h_1 : \mathbb{R} \rightarrow \mathbb{R}, x \mapsto x^2$   
 $\rightsquigarrow 0$

- $h_2 : \mathbb{R}^2 \rightarrow \mathbb{R}, \begin{pmatrix} x \\ y \end{pmatrix} \mapsto x^2 + y^2$   
 $\rightsquigarrow \begin{pmatrix} 0 \\ 0 \end{pmatrix}$

Find the extreme values of

- $f : \mathbb{R}^2 \rightarrow \mathbb{R}, \begin{pmatrix} x \\ y \end{pmatrix} \mapsto 2xy - 5x^2 + 4x - 2y - 4$
- $g : \mathbb{R}^2 \rightarrow \mathbb{R}, \begin{pmatrix} x \\ y \end{pmatrix} \mapsto e^{xy+x^2+y^2}$  (without checking Hessian)



$f$  and  $g$  plot

## Remember: Machine Power intense

- MNIST: Each image is represented by a  $784 \times 1$  vector
- $\mathcal{H}$  is  $784 \times 784$
- intense, despite symmetry of  $\mathcal{H}$  and using mini-batch

## We will "follow the slope" instead

- Recap:  $\nabla_f(x)$  points towards the direction of greatest increase
- Example:  $f(x) = x^2$

# Gradient Descent: Algorithm

- $L(W) = \frac{1}{D} \sum_d L_d(y_{[d]}, X, W)$
- $\nabla_L(W) = \frac{1}{D} \sum_d \nabla_{L_d}(W)(y_{[d]}, X, W)$
- $W_{t+1} = W_t - \eta \cdot \nabla_L(W_t) \equiv W_t + \Delta W_t$ 
  - $\eta$  ( $= \eta_t$ ) is the learning rate
  - $W_0$  is the initial value

Find the extreme values of

- $f(x, y) = x^2 + y^2$
- Starting values:  $x_0 = y_0 = 3$
- $\eta = 0.01$

Note: You might consider a stopping rule

## Plot & Gradient Descent

Find the extreme values of

- $f(x, y) = x^2 + y^2$
- Starting values:  $x_0 = y_0 = 3$
- $\eta = 0.01$

Note: You might consider a stopping rule

Find the extreme values of

- Learning rate: Perhaps later
- Initial values: Roughly
- Gradient: Numerically vs. analytically
  - Numerically: Slow, approximate & easy
  - Analytically: Fast, exact & difficult

# Backpropagation "Algorithm"

repeat:

- 1 Feed the NN (forward propagation)
- 2 Derive the gradient by backprop ("Backward propagation of errors")
- 3 Optimize



Note: For the sigmoid function  $\sigma(x)$ , it holds that

$$\frac{\partial \sigma(x)}{\partial x} = \sigma(x)(1 - \sigma(x))$$