

Neural Networks: Recap Lecture 1 (14.05.19)

Jan Bauer

jan.bauer@dhbw-mannheim.de

16.05.19

Perceptrone (1957, Rosenblatt)

- several binary inputs produces a binary output.
- in \mathbb{R}^2 , it looks like:

$$\begin{aligned} y = \phi(xW + b) &= \begin{cases} 0 & xW + b \leq \text{threshold} \\ 1 & xW + b > \text{threshold} \end{cases} \\ &= \begin{cases} 0 & x_1 w_1 + x_2 w_2 + b \leq \text{threshold} \\ 1 & x_1 w_1 + x_2 w_2 + b > \text{threshold} \end{cases} \end{aligned}$$

Logical operations: (Remember: TRUE = 1, FALSE = 0)

- AND?
 - (TRUE, TRUE) \rightarrow TRUE. (TRUE, FALSE) \rightarrow FALSE...
- OR?
 - (TRUE, TRUE) \rightarrow TRUE. (TRUE, FALSE) \rightarrow TRUE...
- XOR? (Exclusive Or)
 - (TRUE, TRUE) \rightarrow FALSE. (TRUE, FALSE) \rightarrow TRUE...

Solution:

- Nonlinear functions
- More layer (Multilayer Perception) EXAMPLE

Neural Network Architecture [Sketch]

A NN associates to an input X an output $y \equiv f(X, W)$,

- $f : \mathbb{R}^{D \times N} \rightarrow \mathbb{R}^D$
- $X = \begin{pmatrix} x^{(11)} & x^{(12)} & \dots & x^{(1N)} \\ x^{(21)} & x^{(22)} & \dots & x^{(2N)} \\ \vdots & \vdots & \ddots & \vdots \\ x^{(D1)} & x^{(D2)} & \dots & x^{(DN)} \end{pmatrix} = \begin{pmatrix} x^{(1)} \\ x^{(2)} \\ \vdots \\ x^{(D)} \end{pmatrix}$ is $D \times N$
- $W = (w^{(1)} \quad w^{(2)} \quad \dots \quad w^{(N)})'$ is $N \times 1$
- $D = \# \text{data}$ (i.e. training data)

Example

- feed 5 images ($D = 5$), with 784 pixels per image $N = 784$
- Output: 5-dimensional ($D \times 1$) vector, containing the 'guess' of the NN

- Possible Input:
$$\begin{pmatrix} \text{Image of digit 3} \\ \text{Image of digit 0} \\ \text{Image of digit 2} \\ \text{Image of digit 1} \\ \text{Image of digit 9} \end{pmatrix}$$

- Possible output:
$$\begin{pmatrix} \text{"it's a 3!"} \\ \text{"it's a 0!"} \\ \text{"it's a 3!"} \\ \text{"it's a 1!"} \\ \text{"it's a 9!"} \end{pmatrix}$$

Neural Network Architecture [Sketch]

- We will denote the output for the data point d as $y[d]$
- $y[d]$ gives us a score to which class the data fits best
 - TRICKY: Score (vector) or scalar?
 - Score for training ("80
 - Scalar for application ("it's an 8!")

Neuron [Sketch]

The j -th Neuron in layer i associates to an input $Z_{i,j}$ from Neuron k in layer $i - 1$ an output $y_{i,j} \equiv f_{i,j}(f_{i-1,k}, W_{i,j})$, also called scores, s.t.

- $f_{i,j} : \mathbb{R}^{D \times N_{i-1}} \rightarrow \mathbb{R}^{D \times N_i}$
- $Z_{i,j} = f_{i-1,k} W_{i,j} + b_{i,j}$
- $f_{i-1,k}$ is $D \times N_{i-1}$
- $W_{i,j}$ is $N_{i-1} \times N_i$

- Yields the predicted probability for the \hat{k} -th class, given a sample x and a weighting W

$$P(y_{[d]} = \hat{k} | x) = \frac{\exp(y_{[d]}^{(\hat{k})})}{\sum_{k=1}^K \exp(y_{[d]}^{(k)})}$$

- Example: $y_{[d]} = \begin{pmatrix} 5 \\ 0 \\ 3 \\ 2 \end{pmatrix} \rightsquigarrow \begin{pmatrix} 0.8390 \\ 0.0056 \\ 0.1135 \\ 0.0417 \end{pmatrix}$

Sample average over the data loss using a loss function:

$$L(W) \equiv L \equiv \hat{E} L_d \equiv \frac{1}{D} \sum_{d=1}^D L_d(y_{[d]}, X, W)$$

- Loss is high when doing a poor job
- Loss is low when doing a good job

↪ interested in distance between prediction and what is in fact correct

Distance? Just use a norm!

- \mathcal{L}_1 norm

$$L_d = \left\| y_{[d]} - y_{[d]}^* \right\|_1 \equiv \sum_k \left| y_{[d]}^{(k)} - y_{[d]}^{*(k)} \right|$$

- \mathcal{L}_2 norm

$$L_d = \left\| y_{[d]} - y_{[d]}^* \right\|_2^2 \equiv \sum_k \left(y_{[d]}^{(k)} - y_{[d]}^{*(k)} \right)^2$$

Let $x = (x_1, \dots, x_n)$. Then

- $\|x\|_1 \equiv \sum_{i=1}^n |x_i|$

Example: $\|(1, -3)'\|_1 = |1| + |-3| = 1 + 3 = 4$

- $\|x\|_2 \equiv \sqrt{\sum_{i=1}^n x_i^2}$

Example: $\|(1, -3)'\|_2 = \sqrt{1 + 9} = \sqrt{(1)^2 + (-3)^2} = \sqrt{10}$

Note: We might consider to square the \mathcal{L}_2 norm

Example:

Let $y_{[d]} = (1, 2, 3, 4)'$, $y_{[d]}^* = (0, 1, 4, 3)'$

- \mathcal{L}_1 norm

$$L_d = \left\| y_{[d]} - y_{[d]}^* \right\|_1 \rightsquigarrow 4$$

- \mathcal{L}_2 norm (squared)

$$L_d = \left\| y_{[d]} - y_{[d]}^* \right\|_2^2 \rightsquigarrow 4$$

Exercise time:

Consider you feed your NN with three training data points, each of dimension three. As an outcome, you receive

$$y_{[1]} = (0, 1, 3)', y_{[2]} = (1, -1, 5)', y_{[3]} = (2, 1, 0)' .$$

The true values are

$$y_{[1]}^* = (-2, 1, 4)', y_{[2]}^* = (2, 0, 6)', y_{[3]}^* = (0, 0, 0)' .$$

What is the loss, using the \mathcal{L}_1 (the squared \mathcal{L}_2) norm?

Foreshadowing:

- Why is squaring the norm possible?
→ Squaring is a monotone operation
- Why do we even square?
→ Gradient becomes much simpler

plot x^2 and $(x^2)^2 = x^4$